

AUDIT-SID: Artifact-Level Diagnostics for Semantic-ID Tokenizers in Generative Recommendation

Anonymous Author(s)

Abstract

Semantic-ID tokenizers have become an interface layer in generative recommendation: they turn items into discrete addresses that a sequence model can generate. Aggregate Recall@K or NDCG, however, often gives little information about the address space itself. Codebook collapse, full-code collisions, behaviorally weak prefixes, tail compression, and large prefix fan-out can remain hidden until after expensive generator training. We present AUDIT-SID, a mapping-first diagnostic toolkit for item-to-SID tokenizer artifacts. AUDIT-SID normalizes item mappings, metadata, interactions, and optional generator outputs, then reports utilization, collision profile, semantic-collaborative alignment, head-tail allocation, and structural cost proxies. On a same-item Musical case study with 23,742 items, a GRID/RQ-KMeans-style feature-text export has 3,749 unique full SIDs and a 0.9769 full-collision rate, while a bounded ReSID/GAOQ export is collision-free under its exported mapping; the GRID collision rate remains 0.9751–0.9769 across three seeds. Controlled stressors further show that raw collision volume, interaction-qualified collision risk, tail capacity, and active-prefix cost are distinct artifact signals. The contribution is a reusable artifact audit suite, not a new tokenizer or a generative-recommender leaderboard.

CCS Concepts

• **Information systems** → **Recommender systems**; *Retrieval models and ranking*.

Keywords

semantic IDs, generative recommendation, tokenizer diagnostics, recommender systems

ACM Reference Format:

Anonymous Author(s). 2026. AUDIT-SID: Artifact-Level Diagnostics for Semantic-ID Tokenizers in Generative Recommendation. In . ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 Introduction

Semantic-ID (SID) tokenization has become a practical design point for generative recommendation. Systems such as TIGER map items into discrete code sequences so that a sequence model can generate candidate identifiers rather than score every catalog item directly [19]. This shift has made the tokenizer itself a first-order

artifact. Recent work has expanded the space quickly: residual quantization frameworks, recommender-native encoding, collaborative tokenization, differentiable tokenization, non-uniform capacity, qualified collisions, variable-length interfaces, and drift-aware refreshes all change the item address space exposed to a generator [3, 6–8, 10, 14, 22, 25].

The evaluation practice has not caught up with this artifact boundary. Prior recommender tooling has argued for behavioral tests beyond NDCG-style aggregates [4], but SID tokenizers introduce a distinct artifact: the generated item address space. A tokenizer can support a usable downstream Recall@K while still hiding failure modes that are hard to diagnose from aggregate ranking metrics: underused codebooks, repeated full codes, semantically neat but behaviorally poor prefixes, tail items sharing too little capacity, or prefix structures that increase serving cost. These issues matter because SIDs are not merely features. They define the address space exposed to the generator.

This gap is especially awkward for resource reuse. A new tokenizer repository may expose checkpoints, item mappings, codebooks, or only intermediate feature files. Downstream users then need to answer basic engineering questions before training a generator: do all catalog items have valid codes, are many items aliased to the same full code, are collisions concentrated in the tail, and do prefix neighborhoods reflect user co-occurrence rather than only metadata taxonomy? Today these checks are usually reimplemented ad hoc inside each paper. The resulting evidence is difficult to compare because each method reports a different mixture of codebook usage, downstream ranking, and qualitative examples.

We present AUDIT-SID, a resource for auditing item-to-SID tokenizer artifacts before or alongside downstream recommendation evaluation. The goal is deliberately narrower than proposing a new tokenizer. AUDIT-SID asks a concrete question: given a public tokenizer artifact that maps items to code sequences, what can be said about its capacity use, collisions, behavioral alignment, head-tail allocation, and deployment cost without retraining a full generator?

The resource contribution is threefold. First, AUDIT-SID defines a small mapping-first interface for SID assignments, item metadata, interaction histories, and optional generator outputs. Second, it implements D1–D5a diagnostics for utilization, collisions, semantic-collaborative alignment, head-tail capacity, and structural serving-cost proxies, with D6 churn support for continual-tokenizer settings. Third, it demonstrates the audit on public GRID/RQ-KMeans-style, ReSID/GAOQ, sanity, and controlled-stressor artifacts. The same-item Musical case study and stressor table support diagnostic claims: they separate full-code collision pressure, interaction-qualified collision risk, collaborative-prefix recovery, tail capacity, and structural prefix fan-out before a downstream generator is trained. The paper follows the CIKM Resource four-page constraint, where appendices count toward the limit and detailed logs are therefore placed in the online artifact [5].

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference'17, Washington, DC, USA

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-x-xxxx-xxxx-x/YYYY/MM

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

AUDIT-SID: SID exports to diagnostic evidence

Mapping-first checks preview artifact failures before generator training.




Figure 1: AUDIT-SID normalizes heterogeneous SID exports through a mapping-first contract and reports D1–D5a diagnostics. The bottom panels preview representative outputs from the Musical case study and controllers: capacity pressure, full-code collisions, collaborative-prefix recovery, head-tail allocation, and realized prefix cost. D6 is optional churn evidence; D7 requires generator outputs or beam traces.

2 Resource Scope and Interface

AUDIT-SID treats a tokenizer export as an inspectable resource artifact, not as an opaque component of a trained recommender. The minimum input is a table `sid_assignments(item_id, sid_0, ..., sid_L)`. Optional tables provide metadata, interactions, refresh pairs, and generator outputs. This mapping-first contract lets adapters normalize heterogeneous repositories while keeping the audit independent of a particular training loop. Figure 1 summarizes the workflow and previews the diagnostic outputs that the resource is designed to expose.

Artifact contract. The required contract has three checks: each row has a stable item key, every SID level is a discrete token, and the diagnostic item universe is joinable to metadata and interaction slices. Metadata enables semantic/category slices; interactions enable collaborative-neighborhood and head-tail diagnostics. The `generator_outputs` table is optional because many public tokenizer releases expose item codes without trained generator traces.

Evidence roles. AUDIT-SID separates rows by what they can support. A named-method row comes from a public method path and can support method-facing artifact claims. A controlled stressor is deliberately synthetic or method-inspired and can test whether a diagnostic reacts to a mechanism, but it is not method coverage. Optional rows exercise extensions such as drift or non-Amazon

schemas. Literature-only rows motivate the design space but are not loaded as runnable evidence.

Validation before metrics. Adapters are refused as evidence until they pass item-count agreement, missing-code, duplicate-key, depth-consistency, and provenance checks. This guard is necessary because public SID repositories may expose checkpoints, item mappings, codebooks, intermediate feature files, or only partial releases. The manifest records each row’s evidence role and caveats before any D1–D7 number is emitted.

Prior recommender resources such as RecList and Elliot motivate diagnostic and reproducible evaluation beyond a single aggregate metric [1, 4]. AUDIT-SID applies that resource logic to a newer artifact boundary: the item address space consumed by generative recommenders.

3 Diagnostic Design and Method Space

SID papers now optimize different artifact properties: learned item indexing and TIGER-style generative retrieval established the interface [9, 19], while later work stresses collaborative alignment, end-to-end or differentiable tokenization, representation-source sensitivity, code-assignment diversity, collision qualification, variable-length interfaces, drift, and deployment surfaces [2, 3, 7, 8, 11, 15, 16, 18, 21, 24, 25]. The literature space is broader than the artifacts that v0 can honestly execute. Table 1 therefore reports only rows that pass the validator or are explicitly scoped as controls, stressors, or backlog. The method taxonomy guides future adapters; it is not presented as reproduced method coverage.

D1–D4. D1 reports per-level usage, prefix counts, imbalance summaries, and dead or underused code indicators. D2 measures duplicate full SIDs and prefix collisions; it is a collision profile, not a causal harm estimate. A bounded interaction-qualified controller is included because collision-aware work argues that collisions are not uniformly harmful [8]. D3 asks whether prefix neighborhoods recover item co-occurrence neighbors, keeping metadata purity as auxiliary context rather than as collaborative quality. D4 splits items by popularity and measures whether full-code capacity and prefix structure are allocated differently across head, mid, and tail items.

D5a–D7. D5a reports structural cost proxies: SID length, prefix fan-out, duplicate codes, and trie-like expansion pressure. These are not serving latencies; long or variable SID methods make this distinction important [3, 23]. D6 computes churn between tokenizer refreshes for drift-aware settings [2, 6]. D7 is an interface hook for generator behavior, such as invalid paths, next-token entropy, and duplicate generated candidates; it requires `generator_outputs` and is not claimed by the current evidence.

4 Demonstration and Findings

The demonstration evaluates the resource interface rather than ranking tokenizer systems. It combines two public SID export paths with controlled rows that calibrate the diagnostics. GRID supplies the open RQ-KMeans-style path; ReSID supplies the GAOQ line and processed Musical artifacts. Sanity tokenizers and stressors are included only when they clarify what a diagnostic can and cannot identify.

Table 1: Evidence map for the v0 artifact. Literature method-space coverage is discussed in the text; this table reports only artifacts that pass the v0 validator, controlled rows used for diagnostic calibration, and backlog rows that do not support empirical claims.

Row	Cluster	Status	Items	Seeds	Diagnostics	Caveat
GRID/RQ-KMeans feature-text	A	real export	23,742 / 50,000	3 / 1	D1–D5a	processed feature text; not raw-text TIGER/GRID
ReSID/GAOQ bounded	B1/B2	real export	23,742	1	D1–D5a	one-epoch bounded Musical mapping
Category-prefix	control	built-in	23,742	n/a	D1–D5a	deterministic sanity row
Mod-collision hash	control	built-in	23,742	n/a	D1–D5a	collision lower-bound control
Popularity-balanced	control	built-in	23,742	n/a	D1–D5a	popularity calibration row
Qualified collision	stressor	built-in	synthetic	n/a	D2b,D3	not named-method evidence
Capacity budget	stressor	built-in	synthetic	n/a	D1,D2,D4,D5a	not named-method evidence
Variable depth	stressor	built-in	synthetic	n/a	D5a	optional boundary evidence
15+ named SID lines	backlog	not run	–	–	–	adapters require official artifacts

Table 2: Same-item Musical artifact profiles on 23,742 items. Higher unique full codes, D3 L1 co-occurrence prefix recall, and D4 tail unique-SID ratio are better; lower D2 full-collision rate is better. D5a prefixes are level-wise prefix counts, not latency. A dagger marks structural one-full-code-per-item rows where D2=0 and D4 tail=1.0 are forced by the export design.

System	Unique	D2 full	D3 L1	D4 tail	D5a prefixes
GRID ft (3s)	3,857±112	.9759±.0009	.0519±.0037	.377±.008	64/3.4–3.7k/3.7–4.0k
ReSID[†]	23,742	.0000	.1535	1.0000	32/1280/23.7k
<i>Cat-prefix[†]</i>	23,742	.0000	.4470	1.0000	30/83/313/23.7k
<i>Hash-collide</i>	256	1.0000	.0037	.0322	256/256/256/256
<i>Pop-balanced</i>	22,707	.0860	.3026	.9619	4/1024/22.7k

Case-study protocol. The paper-facing case study uses the same 23,742 Musical items for both rows. The GRID row is a controlled feature-text export: ReSID processed feature text is embedded and passed through the official GRID-style residual k-means module. The ReSID row is a bounded FAMA-to-GAOQ export on the same item universe. This modest protocol gives reviewers a same-item diagnostic contrast while avoiding the stronger, unsupported claim that the paper reproduces every training choice of TIGER, GRID, or the full ReSID pipeline.

Table 2 compares two exports on the same item set. The GRID row is intentionally labeled as feature-text: it uses processed feature text from the ReSID artifact path and the official residual quantization module, not a faithful raw-text TIGER reproduction. Under that controlled setup, the row exposes high collision pressure and limited tail capacity. The ReSID/GAOQ row exports one full code per item in this bounded run, yielding zero full-code collisions and higher D3-L1 collaborative prefix recovery. The daggered rows make the structural floor explicit: collision-free full codes are not by themselves evidence of better recommendation quality. Across seeds 42–44, the GRID duplicate-SID rate remains 0.8327–0.8421 and full-collision rate remains 0.9751–0.9769, so the collision pressure is not a single-seed artifact. The comparison is also architecturally biased: GRID is capacity-constrained, while the bounded ReSID export reaches an item-unique leaf level. A same-capacity ablation is future work, not a hidden claim.

Table 3: Controlled stressors isolate diagnostic mechanisms. Each row activates one failure mode and asks whether the corresponding diagnostic separates the two conditions. Full numeric tables are shipped in the artifact.

Stressor	D	Test	Baseline	Under stress
Qualified collision	D2b	co-occur ≠ hash collisions	hash: 1.19x risk	co-occur: 3.86x risk
Capacity budget	D1/D4	tight codebooks skew capacity	head: 1.000 unique	tail: 0.028 unique
Variable depth	D5a	active cost ≠ max-depth view	max-depth: 12,010	active: 7,914

Diagnostic findings. Tables 2 and 3 support three artifact-level findings. First, collision pressure is stable in the controlled GRID row, while the same interface ingests a collision-free ReSID mapping. The qualified-collision stressor then shows why raw collision volume is insufficient: collisions that overlap user co-occurrence have a 3.86x lift, while synthetic hash collisions have 1.19x. Second, collision-free addresses and collaborative-prefix alignment are different objectives. ReSID has no full collisions and higher D3-L1 than the GRID row, while a category-prefix control in the repository recovers more level-1 co-occurrence neighbors without being a learned tokenizer. Third, D4 and D5a separate capacity allocation from prefix structure: in the capacity stressor, head uniqueness reaches 1.000 while tail uniqueness remains 0.028, and variable-depth stressors show that active prefix cost (7,914 realized prefixes) need not match a fixed maximum-depth view (12,010 prefixes).

Repository tables extend these checks without changing the paper claims: sanity rows calibrate D1–D5a, GRID scale rows test larger item counts, DACT rows exercise optional D6 churn, and MovieLens rows check non-Amazon schema portability. Table 2 remains the only method-facing case study in the main paper.

5 Availability, Reproducibility, and Limits

The artifact repository contains the AUDIT-SID Python package, adapter scripts, metric runners, generated CSV/Markdown/L^AT_EX tables, provenance logs, BibTeX/source-audit notes, an MIT license, and a reviewer quickstart. The reviewer entry point is the tag `audit-sid-cikm-resource-v0.1`, with a clean-checkout verifier for the published tables and figure. Large local artifacts are kept out of git and described through manifests so reviewers can separate public reproducibility assets from local caches. The artifact

ships with a clean-checkout verifier script that checks table files, figure inputs, and claim-facing CSVs from manifest-pinned assets; reviewer entry points are documented in the repository quickstart.

The resource is intended to be used in three modes. A paper reviewer can run the clean-check verifier and compare published CSVs with the PDF tables. A method author can add an adapter that emits the normalized `sid_assignments` table, then receive D1–D5a reports without changing their training loop. A downstream recommender researcher can use the same diagnostics as a pre-training gate: artifacts with missing items, inconsistent depth, severe collisions, or weak collaborative-prefix recovery can be flagged before generator training consumes GPU time. These modes are deliberately lighter than a full benchmark suite, but they match the resource-track goal of making a reusable software artifact available at review time.

The evidence supports a conservative resource-demo claim. The toolkit does not reproduce TIGER end-to-end; the GRID Musical row is a controlled feature-text export. The ReSID Sports balanced GAOQ path was stopped after constrained k-means became CPU-bound, so the main ReSID evidence is the bounded Musical export. CARD remains a fallback/stressor path, not faithful CARD reproduction [22]. D2 includes an interaction-qualified controller but not causal harm; D3 is not proven monotonic with Recall@K or NDCG; D5a is structural rather than measured latency; and D6 is optional drift evidence.

AUDIT-SID audits tokenizer artifacts; it is not a new SID tokenizer, a complete generative-recommender evaluation, or an industrial online-impact study. Artifact diagnostics can reveal pressure points before generator training, but they do not replace ranking evaluation. Future versions should add causal collision-harm validation, generator-output diagnostics, stronger same-dataset method coverage, and unified search/recommendation checks [13, 17]. Industrial SID deployment and ranking studies further motivate treating these artifact properties as engineering objects [11, 12, 20].

New adapters emit the same normalized tables, declare their evidence role, and record feature or checkpoint provenance. Versioned manifests and fixed-name “latest” reports preserve stable entry points and timestamped provenance. Thus the Musical case study is not a ranking of GRID against ReSID; it shows how one artifact contract turns collisions, collaborative prefix recovery, tail allocation, and prefix fan-out into separately checkable review objects.

The extension path is also explicit. A new named tokenizer improves method coverage only when it exports item-level codes that pass the validator and can be joined to metadata or interactions. A controlled stressor improves metric calibration but does not count as method coverage. A generator trace activates D7, but without such traces AUDIT-SID reports only mapping-level and diagnostic claims. This separation is what keeps the artifact useful as the SID literature changes: new methods can be added without rewriting the diagnostics, while claims remain tied to the evidence role of each row.

GenAI Usage Disclosure

The authors used generative AI tools for drafting assistance, code-review support, and experiment-log organization. The reported

claims, tables, and citations were checked against the artifact repository and cited sources during paper preparation.

References

- [1] Vito Walter Anelli, Alejandro Bellogin, Antonio Ferrara, Daniele Malatesta, Felice Antonio Merri, Claudio Pomo, Francesco M. Donini, and Tommaso Di Noia. 2021. Elliot: A Comprehensive and Rigorous Framework for Reproducible Recommender Systems Evaluation. In *Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval*. 2405–2414. doi:10.1145/3404835.3463245
- [2] Vladimir Baikalov, Iskander Bagautdinov, and Sergey Muravyov. 2026. Mitigating Collaborative Semantic ID Staleness in Generative Retrieval. *arXiv preprint arXiv:2604.13273* (2026). arXiv:2604.13273 [cs.IR] doi:10.48550/arXiv:2604.13273
- [3] Wenzhuo Cheng, Menghang Gong, Qixin Guo, Hang Zheng, Zhaobin Yang, Jianguo Lou, and Zhengwei Zheng. 2026. CapsID: Soft-Routed Variable-Length Semantic IDs for Generative Recommendation. *arXiv preprint arXiv:2605.05096* (2026). arXiv:2605.05096 [cs.IR] doi:10.48550/arXiv:2605.05096
- [4] Patrick John Chia, Maurizio Ferrari Dacrema, Maria Soledad Pera, and Jacopo Tagliabue. 2022. Beyond NDCG: Behavioral Testing of Recommender Systems with RecList. In *Companion Proceedings of the ACM Web Conference 2022*. arXiv:2111.09963 [cs.IR] doi:10.1145/3487553.3524215 arXiv:2111.09963.
- [5] CIKM 2026 Organizing Committee. 2026. CIKM 2026 Resource Papers. <https://cikm2026.diag.uniroma1.it/resource-papers/>. Accessed 2026-05-19.
- [6] Yuebo Feng, Jiahao Liu, Mingzhe Han, Dongsheng Li, Hansu Gu, Peng Zhang, Tun Lu, and Ning Gu. 2026. Drift-Aware Continual Tokenization for Generative Recommendation. *arXiv preprint arXiv:2603.29705* (2026). arXiv:2603.29705 [cs.IR] doi:10.48550/arXiv:2603.29705
- [7] Junchen Fu, Xuri Ge, Alexandros Karatzoglou, Ioannis Arapakis, Suzan Verberne, Joemon M. Jose, and Zhaochun Ren. 2026. Differentiable Semantic ID for Generative Recommendation. In *Proceedings of the 49th International ACM SIGIR Conference on Research and Development in Information Retrieval*. arXiv:2601.19711 [cs.IR] doi:10.48550/arXiv:2601.19711 Accepted by SIGIR 2026; arXiv:2601.19711.
- [8] Zheng Hu, Yuxin Chen, Yongsan Pan, Xu Yuan, Yuting Yin, Daoyuan Wang, Boyang Xia, Zefei Luo, Hongyang Wang, Songhao Ni, Dongxu Liang, Jun Wang, Shimin Cai, Tao Zhou, Fuji Ren, and Wenwu Ou. 2026. Stop Treating Collisions Equally: Qualification-Aware Semantic ID Learning for Recommendation at Industrial Scale. *arXiv preprint arXiv:2603.00632* (2026). arXiv:2603.00632 [cs.IR] doi:10.48550/arXiv:2603.00632
- [9] Wenyue Hua, Yingqiang Ge, Shuyuan Xu, Zelong Ji, Jiayue Chen, and Yongfeng Zhang. 2023. How to Index Item IDs for Recommendation Foundation Models. *arXiv preprint arXiv:2305.06569* (2023). arXiv:2305.06569 [cs.IR] doi:10.48550/arXiv:2305.06569
- [10] Clark Mingxuan Ju, Liam Collins, Leonardo Neves, Bhuvish Kumar, Louis Yufeng Wang, Tong Zhao, and Neil Shah. 2025. Generative Recommendation with Semantic IDs: A Practitioner’s Handbook. *arXiv preprint arXiv:2507.22224* (2025). arXiv:2507.22224 [cs.IR] doi:10.48550/arXiv:2507.22224
- [11] Clark Mingxuan Ju, Andy Lin, Adedamola Oyewole, Zeel Patel, Hyunjin Park, Jerome Arje, Zhe Dong, Anuj Bhardwaj, Rahul Rastogi, and Neil Shah. 2026. Semantic IDs for Recommendation Systems. *arXiv preprint arXiv:2604.03949* (2026). arXiv:2604.03949 [cs.IR] doi:10.48550/arXiv:2604.03949
- [12] Guowen Li, Yuepeng Zhang, Shunyu Zhang, Yi Zhang, Xiaozhe Jiang, Yi Wang, and Jingwei Zhuo. 2026. SID-Coord: Coordinating Semantic IDs for ID-based Ranking in Short-Video Search. *arXiv preprint arXiv:2604.10471* (2026). arXiv:2604.10471 [cs.IR] doi:10.48550/arXiv:2604.10471
- [13] Yongqi Li, Xinyu Lin, Wenjie Wang, Fuli Feng, Liang Pang, Wenjie Li, Liqiang Nie, Xiangnan He, and Tat-Seng Chua. 2024. A Survey of Generative Search and Recommendation in the Era of Large Language Models. *arXiv preprint arXiv:2404.16924* (2024). arXiv:2404.16924 [cs.IR] doi:10.48550/arXiv:2404.16924
- [14] Yu Liang, Zhongjin Zhang, Yuxuan Zhu, Kerui Zhang, Zhiluohan Guo, Wenhang Zhou, Zonqi Yang, Kangle Wu, Yabo Ni, Anxiang Zeng, Cong Fu, Jianxin Wang, and Jiazhi Xia. 2026. Rethinking Generative Recommender Tokenizer: Recsys-Native Encoding and Semantic Quantization Beyond LLMs. *arXiv preprint arXiv:2602.02338* (2026). arXiv:2602.02338 [cs.IR] doi:10.48550/arXiv:2602.02338
- [15] Enze Liu, Bowen Zheng, Cheng Ling, Lantao Hu, Han Li, and Wayne Xin Zhao. 2024. Generative Recommender with End-to-End Learnable Item Tokenization. *arXiv preprint arXiv:2409.05546* (2024). arXiv:2409.05546 [cs.IR] doi:10.48550/arXiv:2409.05546
- [16] Yongsan Pan, Yuxin Chen, Zheng Hu, Xu Yuan, Daoyuan Wang, Yuting Yin, Songhao Ni, Hongyang Wang, Jun Wang, Fuji Ren, and Wenwu Ou. 2026. Beyond Static Collision Handling: Adaptive Semantic ID Learning for Multimodal Recommendation at Industrial Scale. *arXiv preprint arXiv:2604.23522* (2026). arXiv:2604.23522 [cs.IR] doi:10.48550/arXiv:2604.23522
- [17] Gustavo Penha, Edoardo D’Amico, Marco De Nadai, Enrico Palumbo, Alexandre Tamborrino, Ali Vardasbi, Max Lefarov, Shawn Lin, Timothy Heath, Francesco Fabbri, and Hugues Bouchard. 2025. Semantic IDs for Joint Generative Search and

- Recommendation. In *Proceedings of the 19th ACM Conference on Recommender Systems*. arXiv:2508.10478 [cs.IR] doi:10.48550/arXiv.2508.10478 Late-Breaking Results track; arXiv:2508.10478.
- [18] Shutong Qiao, Wei Yuan, Tong Chen, Xiangyu Zhao, Quoc Viet Hung Nguyen, and Hongzhi Yin. 2026. When Text-as-Vision Meets Semantic IDs in Generative Recommendation: An Empirical Study. *arXiv preprint arXiv:2601.14697* (2026). arXiv:2601.14697 [cs.IR] doi:10.48550/arXiv.2601.14697
- [19] Shashank Rajput, Nikhil Mehta, Anima Singh, Raghunandan H. Keshavan, Trung Vu, Lukasz Heldt, Lichan Hong, Yi Tay, Vinh Q. Tran, Jonah Samost, Maciej Kula, Ed H. Chi, and Maheswaran Sathiamoorthy. 2023. Recommender Systems with Generative Retrieval. In *Advances in Neural Information Processing Systems*. arXiv:2305.05065 [cs.IR] doi:10.48550/arXiv.2305.05065
- [20] Adit Krishan Singh, Vamsi Meduri, Pranjali Awasthi, Yuanyuan Feng, Ameet Trivedi, Afshin Rostamizadeh, and Sanjiv Kumar. 2023. Better Generalization with Semantic IDs: A Case Study in Ranking for Recommendations. *arXiv preprint arXiv:2306.08121* (2023). arXiv:2306.08121 [cs.IR] doi:10.48550/arXiv.2306.08121
- [21] Bowen Wang, Shizhe Diao, Yan Feng, Baijun He, Chen Li, Xiyuan Wang, Jingang Fu, Haiyang Lyu, and Yonggang Chen. 2024. LETTER: A LLM-Enhanced Two-Tower Recommender System. *arXiv preprint arXiv:2405.07314* (2024). arXiv:2405.07314 [cs.IR] doi:10.48550/arXiv.2405.07314
- [22] Yibiao Wei, Jie Zou, Pengfei Zhang, Xiao Ao, Weikang Guo, Zeyu Ma, and Yang Yang. 2026. CARD: Non-Uniform Quantization of Visual Semantic Unit for Generative Recommendation. *arXiv preprint arXiv:2604.26427* (2026). arXiv:2604.26427 [cs.IR] doi:10.48550/arXiv.2604.26427
- [23] Ming Xia, Zhiqin Zhou, Guoxin Ma, and Dongmin Huang. 2026. Unleash the Potential of Long Semantic IDs for Generative Recommendation. *arXiv preprint arXiv:2602.13573* (2026). arXiv:2602.13573 [cs.IR] doi:10.48550/arXiv.2602.13573
- [24] Bowen Zheng, Yupeng Hou, Hongyu Lu, Yu Chen, Wayne Xin Zhao, Ming Chen, and Ji-Rong Wen. 2023. Adapting Large Language Models by Integrating Collaborative Semantics for Recommendation. *arXiv preprint arXiv:2311.09049* (2023). arXiv:2311.09049 [cs.IR] doi:10.48550/arXiv.2311.09049
- [25] Han Zhu, Hao Liu, Zhiwei Liu, Wenqi Sun, Yingxue Zhang, Xinyang Fu, Weijie Liu, and Hui Xiong. 2024. Collaborative Semantic Tokenization for Generative Retrieval. *arXiv preprint arXiv:2404.14774* (2024). arXiv:2404.14774 [cs.IR] doi:10.48550/arXiv.2404.14774